

▼ CHEM277B Homework 8

Trevor Oldham

▼ (A)

Importing the MNIST Data and printing out relevant dimensions

```
import pandas as pd
import numpy as np
import cv2
import matplotlib.pyplot as plt

training_set, validation_set = pd.read_pickle('mnist.pkl')

#examine the shape of the training set
print(len(training_set[0]))
print(len(training_set[0][0]))
print(len(training_set[0][0][0]))
print(len(validation_set[0]))

#normalized the training data and reshape to shape (ndata, nfeatures)
gray_img_train = np.array(training_set[0]).astype('uint8')
gray_img_train = gray_img_train.reshape(60000, 32, 32)
rgb_img_train = np.zeros((60000, 32, 32, 3))
for i in range(len(gray_img_train[0])):
    rgb_img_train[i] = cv2.cvtColor(gray_img_train[i], cv2.COLOR_GRAY2RGB)
print(gray_img_train.shape)
print(rgb_img_train.shape)
rgb_img_train = np.transpose(rgb_img_train, (0, 3, 1, 2)) # (60000, 3, 32, 32)
normalized_train_set_x = rgb_img_train / 255
print(normalized_train_set_x.shape)

plt.show()

gray_img_test = np.array(validation_set[0]).astype('uint8')
gray_img_test = gray_img_test.reshape(10000, 32, 32)
rgb_img_test = np.zeros((10000, 32, 32, 3))
for i in range(len(gray_img_test[0])):
    rgb_img_test[i] = cv2.cvtColor(gray_img_test[i], cv2.COLOR_GRAY2RGB)
print(gray_img_test.shape)
print(rgb_img_test.shape)
rgb_img_test = np.transpose(rgb_img_test, (0, 3, 1, 2)) # (10000, 3, 32, 32)
normalized_test_set_x = rgb_img_test / 255
print(normalized_test_set_x.shape)
```



```
train_set_y = np.array(training_set[1],
test_set_y = np.array(validation_set[1])

60000
32
32
10000
(60000, 32, 32)
(60000, 32, 32, 3)
(60000, 3, 32, 32)
(10000, 32, 32)
(10000, 32, 32, 3)
(10000, 3, 32, 32)
```

Calculate the size of the output after each convolution sequentially applied to the black and white 32x32 image.

- (i). Convolution Filter size of 2x2, number of filters 33, stride of 2, padding of 0
- (ii). Convolution Filter size of 3x3, number of filters 55, stride of 1, padding of 1
- (iii). Convolution Filter size of 3x3, number of filters 77, stride of 1, padding of 1. Followed by a Max Pooling with filter size of 2x2 and stride 2

Using the equation below:

$$\frac{I - K + 2P}{S} + 1$$

where I is the input dimension (32), K is the size of the filter, P is the size of the padding, and S is the stride.

(i)

$$\frac{32 - 2 + 2(0)}{2} + 1 = 16$$

So in this case we get a 16x16 image for each of the 33 feature maps.

(ii)

$$\frac{16 - 3 + 2(1)}{1} + 1 = 16$$

This case returns a 16x16 image for each of the 55 feature maps.

(iii)

$$\frac{16 - 3 + 2(1)}{1} + 1 = 16$$

Then we apply the max pooling operation with 2x2 filter and stride 2:

$$\frac{16 - 2 + 2(0)}{2} + 1 = 8$$

Which yields an 8x8 image for each of the 77 feature maps.

(B)

The MNIST data set was, in fact, in color (RGB). This means the depth of the input image would be 3. Calculate the dimensionality of the output for the following convolutions sequentially applied to a RGB MNIST input:

- (i). Convolution Filter size of 2x2, number of filters 33, stride of 2, padding of 0
- (ii). Convolution Filter size of 3x3, number of filters 55, stride of 1, padding of 1. Followed by a max pooling layer of kernel size 3x3, stride of 1, padding of 0
- (iii). Convolution Filter size of 3x3, number of filters 77, stride of 1, padding of 1. Followed by a Max Pooling with filter size of 2x2 and stride 2.

(i)

$$\frac{32 - 2 + 2(0)}{2} + 1 = 16$$

This convolution yields a 16x16x3 image for each 33 filters.

(ii)

$$\frac{16 - 3 + 2(1)}{1} + 1 = 16$$

Followed by a max pooling layer with F=3, S = 1, we use the equation below:

$$H2 = \frac{H1 - F}{S} + 1$$

$$H2 = \frac{16 - 3}{1} + 1 = 14$$

The calculation for the width is the same and thus we have a 14x14x3 image for each of the 55 filters.

(iii)

$$14 - 3 + 2(1)$$

$$\frac{17 - 5 + 2(1)}{1} + 1 = 14$$

Followed by a max pooling layer with $F=2$, $S=2$, we use the equation below:

$$H2 = \frac{H1 - F}{S} + 1$$

$$H2 = \frac{14 - 2}{2} + 1 = 7$$

So this layer yields an image of $7 \times 7 \times 3$ for each of the 77 filters.

(C)

```
from functools import wraps
from time import time
```

```
def timing(f):
    @wraps(f)
    def wrap(*args, **kw):
        ts = time()
        result = f(*args, **kw)
        te = time()
        print('func:%r took: %2.4f sec' % (f.__name__, te-ts))
        return result
    return wrap
```

```
from torch.optim import SGD, Adam
import torch.nn.functional as F
import random
from tqdm import tqdm
import math
from sklearn.model_selection import train_test_split, KFold
```

```
def create_chunks(complete_list, chunk_size=None, num_chunks=None):
    """
    Cut a list into multiple chunks, each having chunk_size (the last chunk might be
    smaller)
    """
    chunks = []
    if num_chunks is None:
        num_chunks = math.ceil(len(complete_list) / chunk_size)
    elif chunk_size is None:
        chunk_size = math.ceil(len(complete_list) / num_chunks)
    for i in range(num_chunks):
        chunks.append(complete_list[i * chunk_size: (i + 1) * chunk_size])
```

```
return chunks
```

```
class Trainer():
    def __init__(self, model, optimizer_type, learning_rate, epoch, batch_size, input_transform):
        """ The class for training the model
        model: nn.Module
            A pytorch model
        optimizer_type: 'adam' or 'sgd'
        learning_rate: float
        epoch: int
        batch_size: int
        input_transform: func
            transforming input. Can do reshape here
        """
        self.model = model
        if optimizer_type == "sgd":
            self.optimizer = SGD(model.parameters(), learning_rate, momentum=0.9)
        elif optimizer_type == "adam":
            self.optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

        self.epoch = epoch
        self.batch_size = batch_size
        self.input_transform = input_transform
        self.learning_rate = learning_rate

    @timing
    def train(self, inputs, outputs, val_inputs, val_outputs, early_stop=False, l2=False, silent=False):
        """ train self.model with specified arguments
        inputs: np.array, The shape of input_transform(input) should be (ndata, nfeat)
        outputs: np.array shape (ndata,)
        val_inputs: np.array, The shape of input_transform(val_input) should be (ndata, nfeat)
        val_outputs: np.array shape (ndata,)
        early_stop: bool
        l2: bool
        silent: bool. Controls whether or not to print the train and val error during training
        """
        @return
        a dictionary of arrays with train and val losses and accuracies
        """
        ### convert data to tensor of correct shape and type here ###
        inputs = torch.tensor(inputs, dtype=torch.float32)
        outputs = torch.tensor(outputs, dtype=torch.long)
        val_inputs = torch.tensor(val_inputs, dtype=torch.float32)
        val_outputs = torch.tensor(val_outputs, dtype=torch.long)

        losses = []
        accuracies = []
        val_losses = []
        val_accuracies = []
        weights = self.model.state_dict()
```

```

weights = self.model.state_dict()
lowest_val_loss = np.inf

for n_epoch in tqdm(range(self.epoch), leave=False):
    self.model.train()
    batch_indices = list(range(inputs.shape[0]))
    random.shuffle(batch_indices)
    batch_indices = create_chunks(batch_indices, chunk_size=self.batch_size)
    epoch_loss = 0
    epoch_acc = 0
    for batch in batch_indices:
        batch_importance = len(batch) / len(outputs)
        batch_input = inputs[batch]
        batch_output = outputs[batch]
        ### make prediction and compute loss with loss function of your choice
        batch_predictions = self.model.forward(batch_input)
        loss_func = nn.CrossEntropyLoss()
        loss = loss_func(batch_predictions, batch_output)
        if l2:
            ### Compute the loss with L2 regularization ###
            self.optimizer = torch.optim.Adam(model.parameters(), lr = self.lr)
            loss = loss_func(batch_predictions, batch_output)
        self.optimizer.zero_grad()
        loss.backward()
        self.optimizer.step()
        ### Compute epoch_loss and epoch_acc
    epoch_loss, epoch_acc = self.evaluate(inputs, outputs)
    val_loss, val_acc = self.evaluate(val_inputs, val_outputs, print_acc=False)
    if n_epoch % 10 == 0 and not silent:
        print("Epoch %d/%d - Loss: %.3f - Acc: %.3f" % (n_epoch + 1, self.epoch, epoch_loss, epoch_acc))
        print("          Val_loss: %.3f - Val_acc: %.3f" % (val_loss, val_acc))
    losses.append(epoch_loss.detach().numpy())
    accuracies.append(epoch_acc)
    val_losses.append(val_loss.detach().numpy())
    val_accuracies.append(val_acc)
    if early_stop:
        if val_loss < lowest_val_loss:
            lowest_val_loss = val_loss
            weights = self.model.state_dict()

if early_stop:
    self.model.load_state_dict(weights)

return {"losses": losses, "accuracies": accuracies, "val_losses": val_losses}

def evaluate(self, inputs, outputs, print_acc=False):
    """ evaluate model on provided input and output
    inputs: np.array, The shape of input_transform(input) should be (ndata,nfeat)
    outputs: np.array shape (ndata,)
    print_acc: bool

```

```

@return
losses: float
acc: float
"""

inputs = torch.tensor(inputs, dtype=torch.float32)
outputs = torch.tensor(outputs, dtype=torch.long)

loss_func = nn.CrossEntropyLoss()

pred = self.model.forward(inputs)

losses = loss_func(pred, outputs)
#print("pred = ", pred)
#print("truth = " ,outputs)

sum = 0
for i in range(len(outputs)):
    if outputs[i] == torch.argmax(pred[i]):
        sum += 1
acc = sum / len(outputs)
if print_acc:
    print("Accuracy: %.3f" % acc)
return losses, acc

```

```

from torch import nn
import torch
class ConvMLP(nn.Module):
    def __init__(self):
        super(ConvMLP, self).__init__()
        self.conv = nn.Conv2d(3, 3, kernel_size = 5, stride=1, padding = 2)
        self.fc = nn.Linear(3*32*32, 10)
        self.activation = nn.ReLU()

    def forward(self, x):
        x = self.activation(self.conv(x))
        x = nn.Flatten()(x)
        x = self.activation(self.fc(x))
        return x

```

```

from torchsummary import summary
model = ConvMLP()
summary(model, (3, 32, 32))

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 3, 32, 32]	228
ReLU-2	[-1, 3, 32, 32]	0
Linear-3	[-1, 10]	30.730

```

ReLU-4          [-1, 10]          0
=====
Total params: 30,958
Trainable params: 30,958
Non-trainable params: 0
-----
Input size (MB): 0.01
Forward/backward pass size (MB): 0.05
Params size (MB): 0.12
Estimated Total Size (MB): 0.18
=====

```

```

import matplotlib.pyplot as plt
from sklearn.model_selection import KFold
def Kfold_validation(n, inputs, outputs):

    total_num=len(inputs)
    kf=KFold(n_splits=n,shuffle=True)

    for train_selector,test_selector in kf.split(range(total_num)):
        ### Decide training examples and testing examples for this fold ###
        train_Xs= inputs[train_selector]
        test_Xs= inputs[test_selector]
        train_ys= outputs[train_selector]
        test_ys= outputs[test_selector]

        model = ConvMLP()
        t = Trainer(model, optimizer_type='adam', learning_rate=1e-3, epoch=30, batch_size=16)
        train_in,val_in,train_real,val_real=train_test_split(train_Xs,train_ys, test_Xs,test_ys)

        dictionary = t.train(train_in, train_real, val_in, val_real,early_stop=True)

    return dictionary, model

dictionary, model = Kfold_validation(3, normalized_train_set_x[0:1000], train_set_y)

0%|          | 0/30 [00:00<?, ?it/s]<ipython-input-141-5a654197e45b>:125: U
inputs = torch.tensor(inputs, dtype=torch.float32)
<ipython-input-141-5a654197e45b>:126: UserWarning: To copy construct from a t
outputs = torch.tensor(outputs, dtype=torch.long)
7%|█         | 2/30 [00:00<00:03, 7.84it/s]Epoch 1/30 - Loss: 2.299 - Acc:
Val_loss: 2.301 - Val_acc: 0.103
40%|██████    | 12/30 [00:01<00:02, 7.90it/s]Epoch 11/30 - Loss: 2.282 - Acc:
Val_loss: 2.295 - Val_acc: 0.126
73%|██████████| 22/30 [00:02<00:00, 8.01it/s]Epoch 21/30 - Loss: 2.281 - Acc:
Val_loss: 2.292 - Val_acc: 0.126
func:'train' took: 3.8070 sec
3%|█         | 1/30 [00:00<00:03, 8.41it/s]Epoch 1/30 - Loss: 2.301 - Acc:
Val_loss: 2.305 - Val_acc: 0.090

```



```

40%|██████    | 12/30 [00:01<00:02, 6.09it/s]Epoch 11/30 - Loss: 2.272 - Acc:
Val_loss: 2.314 - Val_acc: 0.094
73%|██████████| 22/30 [00:03<00:01, 6.72it/s]Epoch 21/30 - Loss: 2.271 - Acc:
Val_loss: 2.318 - Val_acc: 0.099
func:'train' took: 4.5425 sec
3%||          | 1/30 [00:00<00:03, 7.63it/s]Epoch 1/30 - Loss: 2.301 - Acc:
Val_loss: 2.303 - Val_acc: 0.090
40%|██████    | 12/30 [00:01<00:02, 7.92it/s]Epoch 11/30 - Loss: 2.281 - Acc:
Val_loss: 2.301 - Val_acc: 0.135
73%|██████████| 22/30 [00:02<00:00, 8.18it/s]Epoch 21/30 - Loss: 2.280 - Acc:
Val_loss: 2.304 - Val_acc: 0.135
func:'train' took: 3.7226 sec

```

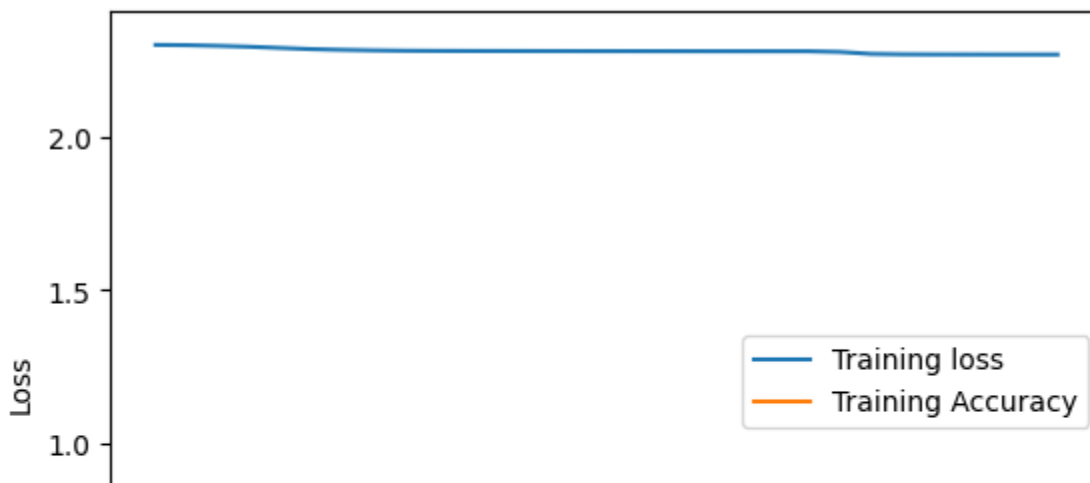
```


losses = dictionary['losses']
accuracies = dictionary['accuracies']
accuracies = np.asarray(accuracies)
losses = np.asarray(losses)
plt.figure()
plt.plot(np.arange(len(losses)),losses,label='Training loss')
plt.plot(np.arange(len(accuracies)),accuracies,label='Training Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

val_losses = dictionary['val_losses']
val_accuracies = dictionary['val_accuracies']
val_accuracies = np.asarray(val_accuracies)
val_losses = np.asarray(val_losses)
plt.figure()
plt.plot(np.arange(len(val_losses)),val_losses,label='Validation loss')
plt.plot(np.arange(len(val_accuracies)),val_accuracies,label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

```

<matplotlib.legend.Legend at 0x7f716c277d00>





```
#feed forward into the model with the test data from MNIST and compare the results
output = model.forward(torch.tensor(normalized_test_set_x, dtype = torch.float32))
sum = 0
for i in range(len(output)):
    if (test_set_y[i] == torch.argmax(output[i])):
        sum += 1
print(sum/len(output))

0.1034
```

(D)

```

class ConvMLP2(nn.Module):
    def __init__(self):
        super(ConvMLP2, self).__init__()
        self.conv = nn.ModuleList([nn.Conv2d(3, 16, kernel_size=3, stride=1, padding=1),
                                    nn.Conv2d(16, 16, kernel_size=3, stride=1, padding=1)])
        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc = nn.ModuleList([nn.Linear(16 * 8 * 8, 128), nn.Linear(128, 10)])
        self.activation = nn.ReLU()

    def forward(self, x):
        x = self.pool(self.activation(self.conv[0](x)))
        x = self.pool(self.activation(self.conv[1](x)))
        x = nn.Flatten()(x)
        x = self.activation(self.fc[0](x))
        x = self.activation(self.fc[1](x))
        return x

model = ConvMLP2()
summary(model, (3, 32, 32))

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 16, 32, 32]	448
ReLU-2	[-1, 16, 32, 32]	0
MaxPool2d-3	[-1, 16, 16, 16]	0
Conv2d-4	[-1, 16, 16, 16]	2,320
ReLU-5	[-1, 16, 16, 16]	0
MaxPool2d-6	[-1, 16, 8, 8]	0
Linear-7	[-1, 128]	131,200
ReLU-8	[-1, 128]	0
Linear-9	[-1, 10]	1,290
ReLU-10	[-1, 10]	0
Total params: 135,258		
Trainable params: 135,258		
Non-trainable params: 0		
Input size (MB): 0.01		
Forward/backward pass size (MB): 0.35		
Params size (MB): 0.52		
Estimated Total Size (MB): 0.88		

```

def Kfold_validation_2(n, inputs, outputs):

    total_num=len(inputs)
    kf=KFold(n_splits=n,shuffle=True)

```

```

for train_selector, test_selector in kf.split(range(total_num)):
    ### Decide training examples and testing examples for this fold ###
    train_Xs= inputs[train_selector]
    test_Xs=  inputs[test_selector]
    train_ys= outputs[train_selector]
    test_ys=  outputs[test_selector]

    model = ConvMLP2()
    t = Trainer(model, optimizer_type='adam', learning_rate=1e-3, epoch=30, batch_size=100,
                train_in, val_in, train_real, val_real=train_test_split(train_Xs, train_ys, test_size=0.2))

    dictionary = t.train(train_in, train_real, val_in, val_real, early_stop=False)

return dictionary, model

```

```
dictionary, model = Kfold_validation_2(3, normalized_train_set_x[0:1000], train_set_y[0:1000])
```

```

0%|          | 0/30 [00:00<?, ?it/s]<ipython-input-141-5a654197e45b>:125: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True) rather than torch.tensor(sourceTensor).
inputs = torch.tensor(inputs, dtype=torch.float32)
<ipython-input-141-5a654197e45b>:126: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True) rather than torch.tensor(sourceTensor).
outputs = torch.tensor(outputs, dtype=torch.long)
3%||         | 1/30 [00:00<00:08, 3.36it/s]Epoch 1/30 - Loss: 2.301 - Acc: 0.099
Val_loss: 2.300 - Val_acc: 0.130
37%|██████   | 11/30 [00:02<00:04, 4.58it/s]Epoch 11/30 - Loss: 2.300 - Acc: 0.112
Val_loss: 2.301 - Val_acc: 0.112
70%|███████   | 21/30 [00:05<00:02, 3.24it/s]Epoch 21/30 - Loss: 2.300 - Acc: 0.112
Val_loss: 2.301 - Val_acc: 0.112
func:'train' took: 7.7381 sec
3%||         | 1/30 [00:00<00:06, 4.73it/s]Epoch 1/30 - Loss: 2.298 - Acc: 0.099
Val_loss: 2.301 - Val_acc: 0.099
37%|██████   | 11/30 [00:02<00:04, 4.64it/s]Epoch 11/30 - Loss: 2.271 - Acc: 0.126
Val_loss: 2.306 - Val_acc: 0.126
70%|███████   | 21/30 [00:04<00:01, 4.64it/s]Epoch 21/30 - Loss: 2.253 - Acc: 0.126
Val_loss: 2.300 - Val_acc: 0.126
func:'train' took: 6.4801 sec
3%||         | 1/30 [00:00<00:06, 4.75it/s]Epoch 1/30 - Loss: 2.301 - Acc: 0.099
Val_loss: 2.304 - Val_acc: 0.090
37%|██████   | 11/30 [00:02<00:04, 3.83it/s]Epoch 11/30 - Loss: 2.281 - Acc: 0.117
Val_loss: 2.308 - Val_acc: 0.117
70%|███████   | 21/30 [00:05<00:02, 3.87it/s]Epoch 21/30 - Loss: 2.255 - Acc: 0.117
Val_loss: 2.304 - Val_acc: 0.117
func:'train' took: 7.4687 sec

```

```

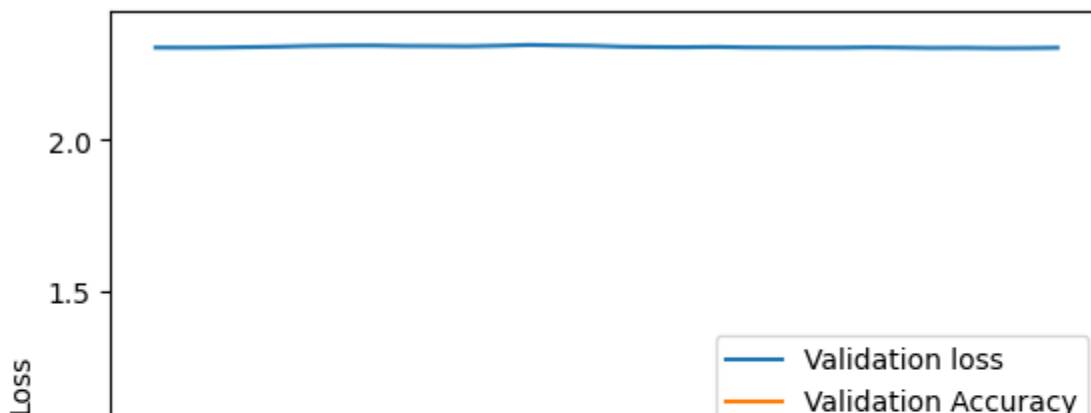

losses = dictionary['losses']
accuracies = dictionary['accuracies']
accuracies = np.asarray(accuracies)
losses = np.asarray(losses)
plt.figure()
plt.plot(np.arange(len(losses)), losses, label='Training loss')

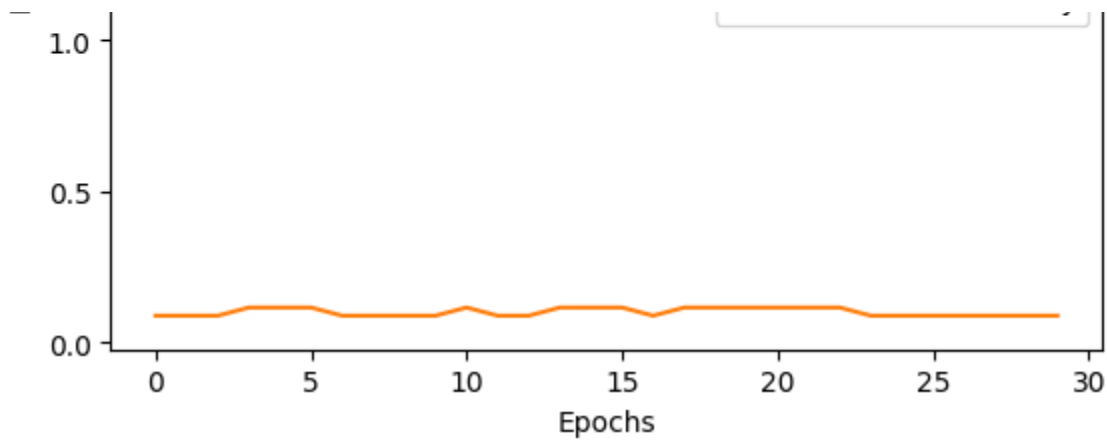
```

```
plt.plot(np.arange(len(accuracies)),accuracies,label='Training Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

val_losses = dictionary['val_losses']
val_accuracies = dictionary['val_accuracies']
val_accuracies = np.asarray(val_accuracies)
val_losses = np.asarray(val_losses)
plt.figure()
plt.plot(np.arange(len(val_losses)),val_losses,label='Validation loss')
plt.plot(np.arange(len(val_accuracies)),val_accuracies,label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
```

<matplotlib.legend.Legend at 0x7f7199390d60>





[Colab paid products](#) - [Cancel contracts here](#)