

▼ CHEM277B Homework 9

Trevor Oldham

(A)

Process the MNIST data set into a training and test set, with one channel of 32x32 images.

```
import pandas as pd
import numpy as np
import cv2
import matplotlib.pyplot as plt

training_set, validation_set = pd.read_pickle('mnist.pkl')

#examine the shape of the training set
print(len(training_set[0]))
print(len(training_set[0][0]))
print(len(training_set[0][0][0]))
print(len(validation_set[0]))

#normalized the training data and reshape to shape (ndata, nfeatures)
gray_img_train = np.array(training_set[0]).astype('uint8')
gray_img_train = gray_img_train.reshape(60000, 32, 32)

normalized_train_set_x = gray_img_train / 255
normalized_train_set_x = normalized_train_set_x.reshape((-1, 1, 32, 32))

print(normalized_train_set_x.shape)

gray_img_test = np.array(validation_set[0]).astype('uint8')
gray_img_test = gray_img_test.reshape(10000, 32, 32)

normalized_test_set_x = gray_img_test / 255
normalized_test_set_x = normalized_test_set_x.reshape((-1, 1, 32, 32))

print(normalized_test_set_x.shape)

train_set_y = np.array(training_set[1])
test_set_y = np.array(validation_set[1])

60000
32
32
10000
(60000, 1, 32, 32)
```



Create a pytorch model based on the specifications in the homework problem.

```
import torch
from torch import nn
class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv = nn.ModuleList([nn.Conv2d(1, 6, kernel_size=6, padding=2), # be1
                                     nn.Conv2d(6, 12, kernel_size=5, padding=2)]) # (E
        self.pooling = nn.ModuleList([nn.MaxPool2d(kernel_size=3), nn.MaxPool2d(ke
        self.fc = nn.ModuleList([nn.Linear(300,300),nn.Linear(300,10)])
        self.activation = nn.ReLU()
        self.bn = [nn.BatchNorm2d(6), nn.BatchNorm2d(12)]

    def forward(self, inp):
        x = self.bn[0](self.conv[0](inp))
        x = self.pooling[0](self.activation(x))
        x = self.pooling[1](self.activation(self.bn[1](self.conv[1](x))))
        x = nn.Flatten()(x)
        residual = x
        y = self.fc[0](x)
        y = y + residual
        y = self.activation(y)
        y = nn.Softmax(dim = -1)(self.fc[1](y))
        return y

from torchsummary import summary
model = CNN()
summary(model, (1, 32, 32))
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 6, 31, 31]	222
ReLU-2	[-1, 6, 31, 31]	0
MaxPool2d-3	[-1, 6, 10, 10]	0
Conv2d-4	[-1, 12, 10, 10]	1,812
ReLU-5	[-1, 12, 10, 10]	0
MaxPool2d-6	[-1, 12, 5, 5]	0
Linear-7	[-1, 300]	90,300
ReLU-8	[-1, 300]	0
Linear-9	[-1, 10]	3,010

```
Total params: 95,344
Trainable params: 95,344
Non-trainable params: 0
```

```
Input size (MB): 0.00
```

```

Forward/backward pass size (MB): 0.12
Params size (MB): 0.36
Estimated Total Size (MB): 0.49
-----

```

Create a trainer class to train the model.

```

from functools import wraps
from time import time

def timing(f):
    @wraps(f)
    def wrap(*args, **kw):
        ts = time()
        result = f(*args, **kw)
        te = time()
        print('func:%r    took: %2.4f sec' % (f.__name__, te-ts))
        return result
    return wrap

from torch.optim import SGD, Adam
import torch.nn.functional as F
import random
from tqdm import tqdm
import math
from sklearn.model_selection import train_test_split, KFold

def create_chunks(complete_list, chunk_size=None, num_chunks=None):
    """
    Cut a list into multiple chunks, each having chunk_size (the last chunk might be
    smaller)
    """
    chunks = []
    if num_chunks is None:
        num_chunks = math.ceil(len(complete_list) / chunk_size)
    elif chunk_size is None:
        chunk_size = math.ceil(len(complete_list) / num_chunks)
    for i in range(num_chunks):
        chunks.append(complete_list[i * chunk_size: (i + 1) * chunk_size])
    return chunks

class Trainer():
    def __init__(self, model, optimizer_type, learning_rate, epoch, batch_size, in):
        """ The class for training the model """
        model: nn.Module
            A pytorch model
        optimizer_type: 'adam' or 'sgd'
        learning_rate: float
        epoch: int
        batch_size: int

```

```

        batch_size: int
        input_transform: func
            transforming input. Can do reshape here
        """

        self.model = model
        if optimizer_type == "sgd":
            self.optimizer = SGD(model.parameters(), learning_rate, momentum=0.9)
        elif optimizer_type == "adam":
            self.optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

        self.epoch = epoch
        self.batch_size = batch_size
        self.input_transform = input_transform
        self.learning_rate = learning_rate

    @timing
    def train(self, inputs, outputs, val_inputs, val_outputs, early_stop=False, l2=False,
            """ train self.model with specified arguments
            inputs: np.array, The shape of input_transform(input) should be (ndata, nfeat)
            outputs: np.array shape (ndata,)
            val_inputs: np.array, The shape of input_transform(val_input) should be (ndata, nfeat)
            val_outputs: np.array shape (ndata,)
            early_stop: bool
            l2: bool
            silent: bool. Controls whether or not to print the train and val error during training

    @return
    a dictionary of arrays with train and val losses and accuracies
    """
        ### convert data to tensor of correct shape and type here ###
        inputs = torch.tensor(inputs, dtype=torch.float32)
        outputs = torch.tensor(outputs, dtype=torch.long)
        val_inputs = torch.tensor(val_inputs, dtype=torch.float32)
        val_outputs = torch.tensor(val_outputs, dtype=torch.long)

        losses = []
        accuracies = []
        val_losses = []
        val_accuracies = []
        weights = self.model.state_dict()
        lowest_val_loss = np.inf

        for n_epoch in tqdm(range(self.epoch), leave=False):
            self.model.train()
            batch_indices = list(range(inputs.shape[0]))
            random.shuffle(batch_indices)
            batch_indices = create_chunks(batch_indices, chunk_size=self.batch_size)
            epoch_loss = 0
            epoch_acc = 0
            for batch in batch_indices:

```

```

        batch_importance = len(batch) / len(outputs)
        batch_input = inputs[batch]
        batch_output = outputs[batch]
        ### make prediction and compute loss with loss function of your choice
        batch_predictions = self.model.forward(batch_input)
        loss_func = nn.CrossEntropyLoss()
        loss = loss_func(batch_predictions, batch_output)
        if l2:
            ### Compute the loss with L2 regularization ###
            self.optimizer = torch.optim.Adam(model.parameters(), lr = self.lr)
            loss = loss_func(batch_predictions, batch_output)
            self.optimizer.zero_grad()
            loss.backward()
            self.optimizer.step()
            ### Compute epoch_loss and epoch_acc
        epoch_loss, epoch_acc = self.evaluate(inputs, outputs)
        val_loss, val_acc = self.evaluate(val_inputs, val_outputs, print_acc=False)
        if n_epoch % 10 == 0 and not silent:
            print("Epoch %d/%d - Loss: %.3f - Acc: %.3f" % (n_epoch + 1, self.loss, val_loss, val_acc))
            print("              Val_loss: %.3f - Val_acc: %.3f" % (val_loss, val_acc))
        losses.append(epoch_loss.detach().numpy())
        accuracies.append(epoch_acc)
        val_losses.append(val_loss.detach().numpy())
        val_accuracies.append(val_acc)
        if early_stop:
            if val_loss < lowest_val_loss:
                lowest_val_loss = val_loss
                weights = self.model.state_dict()

    if early_stop:
        self.model.load_state_dict(weights)

    return {"losses": losses, "accuracies": accuracies, "val_losses": val_losses, "val_accs": val_accs}

def evaluate(self, inputs, outputs, print_acc=False):
    """ evaluate model on provided input and output
    inputs: np.array, The shape of input_transform(input) should be (ndata,nfeatures)
    outputs: np.array shape (ndata,)
    print_acc: bool

    @return
    losses: float
    acc: float
    """

    inputs = torch.tensor(inputs, dtype=torch.float32)
    outputs = torch.tensor(outputs, dtype=torch.long)

    loss_func = nn.CrossEntropyLoss()

    npred = self.model.forward(inputs)

```

```

pred = net.module.forward(inputs,

losses = loss_func(pred, outputs)
#print("pred = ", pred)
#print("truth = " ,outputs)

sum = 0
for i in range(len(outputs)):
    if outputs[i] == torch.argmax(pred[i]):
        sum += 1
acc = sum / len(outputs)
if print_acc:
    print("Accuracy: %.3f" % acc)
return losses, acc

```

Run the model with and without batch normalization. Which give you better test accuracy?

```

import matplotlib.pyplot as plt

def train_and_validate(model, inputs, outputs):

    total_num=len(inputs)

    model = model
    t = Trainer(model, optimizer_type='adam', learning_rate=1e-3, epoch=30, batch_size=10,
train_in, val_in, train_real, val_real=train_test_split(inputs, outputs, train_size = 0.8))

    dictionary = t.train(train_in, train_real, val_in, val_real, early_stop=True, l2=False)

    return dictionary, model

```

Running the model with batch normalization:

```

model = CNN()
dictionary_1, model_1 = train_and_validate(model, normalized_train_set_x[0:10000],

0%|          | 0/30 [00:00<?, ?it/s]<ipython-input-23-5a654197e45b>:125: Use
inputs = torch.tensor(inputs, dtype=torch.float32)
<ipython-input-23-5a654197e45b>:126: UserWarning: To copy construct from a tensor, use
outputs = torch.tensor(outputs, dtype=torch.long)
3%|          | 1/30 [00:03<01:44, 3.61s/it]Epoch 1/30 - Loss: 1.684 - Acc:
Val_loss: 1.695 - Val_acc: 0.794
37%|██████    | 11/30 [00:43<01:14, 3.94s/it]Epoch 11/30 - Loss: 1.478 - Acc:
Val_loss: 1.494 - Val_acc: 0.971
70%|██████████| 21/30 [01:23<00:37, 4.13s/it]Epoch 21/30 - Loss: 1.469 - Acc:
Val_loss: 1.487 - Val_acc: 0.976
func:'train' took: 121.5386 s

```

Create a new CNN model without batch normalization.

```
class CNN_2(nn.Module):
    def __init__(self):
        super(CNN_2, self).__init__()
        self.conv = nn.ModuleList([nn.Conv2d(1, 6, kernel_size=6, padding=2), # be1
                                     nn.Conv2d(6, 12, kernel_size=5, padding=2)]) # (E
        self.pooling = nn.ModuleList([nn.MaxPool2d(kernel_size=3), nn.MaxPool2d(ke
        self.fc = nn.ModuleList([nn.Linear(300,300),nn.Linear(300,10)])
        self.activation = nn.ReLU()

    def forward(self, inp):
        x = self.conv[0](inp)
        x = self.pooling[0](self.activation(x))
        x = self.pooling[1](self.activation(self.conv[1](x)))
        x = nn.Flatten()(x)
        residual = x
        y = self.fc[0](x)
        y = y + residual
        y = self.activation(y)
        y = nn.Softmax(dim = -1)(self.fc[1](y))
        return y

model = CNN_2()
summary(model, (1, 32, 32))
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 6, 31, 31]	222
ReLU-2	[-1, 6, 31, 31]	0
MaxPool2d-3	[-1, 6, 10, 10]	0
Conv2d-4	[-1, 12, 10, 10]	1,812
ReLU-5	[-1, 12, 10, 10]	0
MaxPool2d-6	[-1, 12, 5, 5]	0
Linear-7	[-1, 300]	90,300
ReLU-8	[-1, 300]	0
Linear-9	[-1, 10]	3,010
Total params: 95,344		
Trainable params: 95,344		
Non-trainable params: 0		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.12		
Params size (MB): 0.36		
Estimated Total Size (MB): 0.49		

```
model = CNN_2()
```

```

model = CNN_2()
dictionary_2, model_2 = train_and_validate(model, normalized_train_set_x[0:10000],

0%|          | 0/30 [00:00<?, ?it/s]<ipython-input-23-5a654197e45b>:125: Use
inputs = torch.tensor(inputs, dtype=torch.float32)
<ipython-input-23-5a654197e45b>:126: UserWarning: To copy construct from a ten
outputs = torch.tensor(outputs, dtype=torch.long)
3%||          | 1/30 [00:03<01:51, 3.84s/it]Epoch 1/30 - Loss: 1.820 - Acc:
Val_loss: 1.807 - Val_acc: 0.693
37%|████      | 11/30 [00:34<00:58, 3.08s/it]Epoch 11/30 - Loss: 1.582 - Acc
Val_loss: 1.592 - Val_acc: 0.873
70%|████████  | 21/30 [01:07<00:28, 3.18s/it]Epoch 21/30 - Loss: 1.487 - Acc
Val_loss: 1.500 - Val_acc: 0.964
func:'train' took: 95.7824 sec

```

Plot the resulting Training Error and Training Accuracy

```

fig, ax = plt.subplots(1, 1, figsize=[6, 4])
losses_1 = dictionary_1['losses']
losses_1 = np.asarray(losses_1)
losses_2 = dictionary_2['losses']
losses_2 = np.asarray(losses_2)

plt.plot(np.arange(len(losses_1)), losses_1, label='With Batch Normalization')
plt.plot(np.arange(len(losses_2)), losses_2, label='Without Batch Normalization')
plt.xlabel('Epochs')
plt.ylabel('Validation Loss')
plt.title("Validation Loss with and without Batch Normalization")
plt.legend()
plt.show()

```

```
#####
```


```

fig, ax = plt.subplots(1, 1, figsize=[6, 4])
acc_1 = dictionary_1['accuracies']
acc_1 = np.asarray(acc_1)
acc_2 = dictionary_2['accuracies']
acc_2 = np.asarray(acc_2)


plt.plot(np.arange(len(acc_1)), acc_1, label='With Batch Normalization')
plt.plot(np.arange(len(acc_2)), acc_2, label='Without Batch Normalization')
plt.xlabel('Epochs')
plt.ylabel('Validation Accuracy')
plt.title("Validation Accuracy with and without Batch Normalization")
plt.legend()

```





<matplotlib.legend.Legend at 0x7f21c82079a0>



The testing accuracy is best when the model is run with batch normalization enabled, and the loss is minimized faster over fewer training epochs. Both models approach a similar value for validation accuracy but the first model trains faster.

(B)

Run the model with and without the skip connection at learning rate of $5e-3$ for 10 epochs. Do you see faster training and/or better test accuracy with the skip connection?

```

### define a model with no skip connection
class CNN_3(nn.Module):
    def __init__(self):
        super(CNN_3, self).__init__()
        self.conv = nn.ModuleList([nn.Conv2d(1, 6, kernel_size=6, padding=2), # be1
                                     nn.Conv2d(6, 12, kernel_size=5, padding=2)]) # (E
        self.pooling = nn.ModuleList([nn.MaxPool2d(kernel_size=3), nn.MaxPool2d(ke
        self.fc = nn.ModuleList([nn.Linear(300,300),nn.Linear(300,10)])
        self.activation = nn.ReLU()
        self.bn = [nn.BatchNorm2d(6), nn.BatchNorm2d(12)]

    def forward(self, inp):
        x = self.bn[0](self.conv[0](inp))
        x = self.pooling[0](self.activation(x))
        x = self.pooling[1](self.activation(self.bn[1](self.conv[1](x))))
        x = nn.Flatten()(x)

        x = self.fc[0](x)

        x = self.activation(x)
        x = nn.Softmax(dim = -1)(self.fc[1](x))
        return x

```

Define a new train and validation function with a learning rate of 5e-3 and 10 epochs.

```

def train_and_validate_2(model, inputs, outputs):

    total_num=len(inputs)

    model = model
    t = Trainer(model, optimizer_type='adam', learning_rate=5e-3, epoch=10, batch_size=10)
    train_in,val_in,train_real,val_real=train_test_split(inputs,outputs, train_size = 0.8)

    dictionary = t.train(train_in, train_real, val_in, val_real,early_stop=True,l2=False)

    return dictionary, model

```

Run the model without the skip connection and record the validation loss and accuracy over 10 epochs.

```

model = CNN_3()
dictionary_3, model_3 = train_and_validate_2(model, normalized_train_set_x[0:10000]

0%|          | 0/10 [00:00<?, ?it/s]<ipython-input-23-5a654197e45b>:125: Use
inputs = torch.tensor(inputs, dtype=torch.float32)
<ipython-input-23-5a654197e45b>:126: UserWarning: To copy construct from a te

```

```

    outputs = torch.tensor(outputs, dtype=torch.long)
10%|█          | 1/10 [00:04<00:37, 4.14s/it]Epoch 1/10 - Loss: 1.599 - Acc:
    Val_loss: 1.610 - Val_acc: 0.854
                                func:'train' took: 40.2127 se

```

Run the model with the skip connection and record the validation loss and accuracy over 10 epochs.

```

model = CNN()
dictionary_4, model_4 = train_and_validate_2(model, normalized_train_set_x[0:10000])

0%|          | 0/10 [00:00<?, ?it/s]<ipython-input-23-5a654197e45b>:125: Use
    inputs = torch.tensor(inputs, dtype=torch.float32)
<ipython-input-23-5a654197e45b>:126: UserWarning: To copy construct from a ten
    outputs = torch.tensor(outputs, dtype=torch.long)
10%|█          | 1/10 [00:03<00:35, 3.99s/it]Epoch 1/10 - Loss: 1.794 - Acc:
    Val_loss: 1.809 - Val_acc: 0.653
                                func:'train' took: 40.2632 se

```

```

fig, ax = plt.subplots(1, 1, figsize=[6, 4])
losses_3 = dictionary_3['losses']
losses_3 = np.asarray(losses_3)
losses_4 = dictionary_4['losses']
losses_4 = np.asarray(losses_4)

plt.plot(np.arange(len(losses_3)), losses_3, label='Without Skip Connection')
plt.plot(np.arange(len(losses_4)), losses_4, label = 'With Skip Connection')
plt.xlabel('Epochs')
plt.ylabel('Validation Loss')
plt.title("Validation Loss with and without Skip Connection")
plt.legend()
plt.show()

```

#####

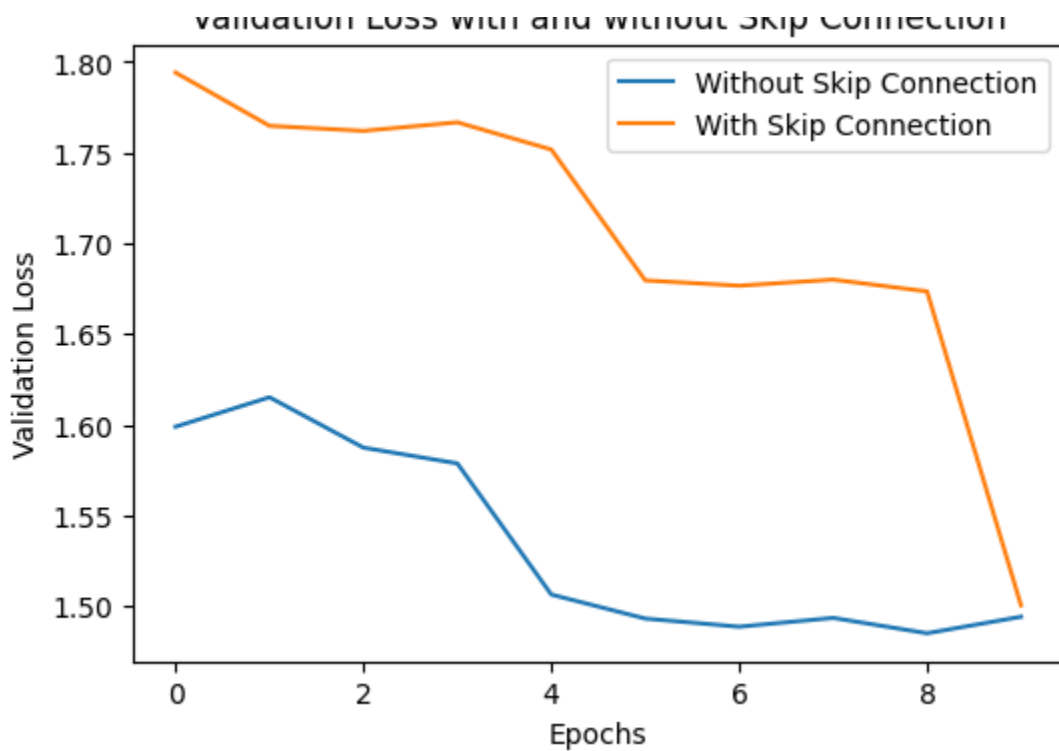
```

fig, ax = plt.subplots(1, 1, figsize=[6, 4])
acc_3 = dictionary_3['accuracies']
acc_3 = np.asarray(acc_3)
acc_4 = dictionary_4['accuracies']
acc_4 = np.asarray(acc_4)


plt.plot(np.arange(len(acc_3)), acc_3, label='Without Skip Connection')
plt.plot(np.arange(len(acc_4)), acc_4, label = 'With Skip Connection')
plt.xlabel('Epochs')
plt.ylabel('Validation Accuracy')
plt.title("Validation Accuracy with and without Batch Normalization")
plt.legend()

```

Validation Loss with and without Skip Connection



<matplotlib.legend.Legend at 0x7f21c87753a0>



It seems counterintuitive but the results show that the model without a skip connection will start out with higher accuracy and lower loss, but at the point of epoch 10 the two models approach the same value. I imagine that the model with the recurrence layer will begin to perform better than the other model as the epochs increase past ten, and eventually improve to a higher accuracy as the training progresses further. We could run the model again with 30 epochs to observe if this is true

observe if this is true.

[Colab paid products](#) - [Cancel contracts here](#)